

Processes and Threads

ECSE 420 - Tutorial 1

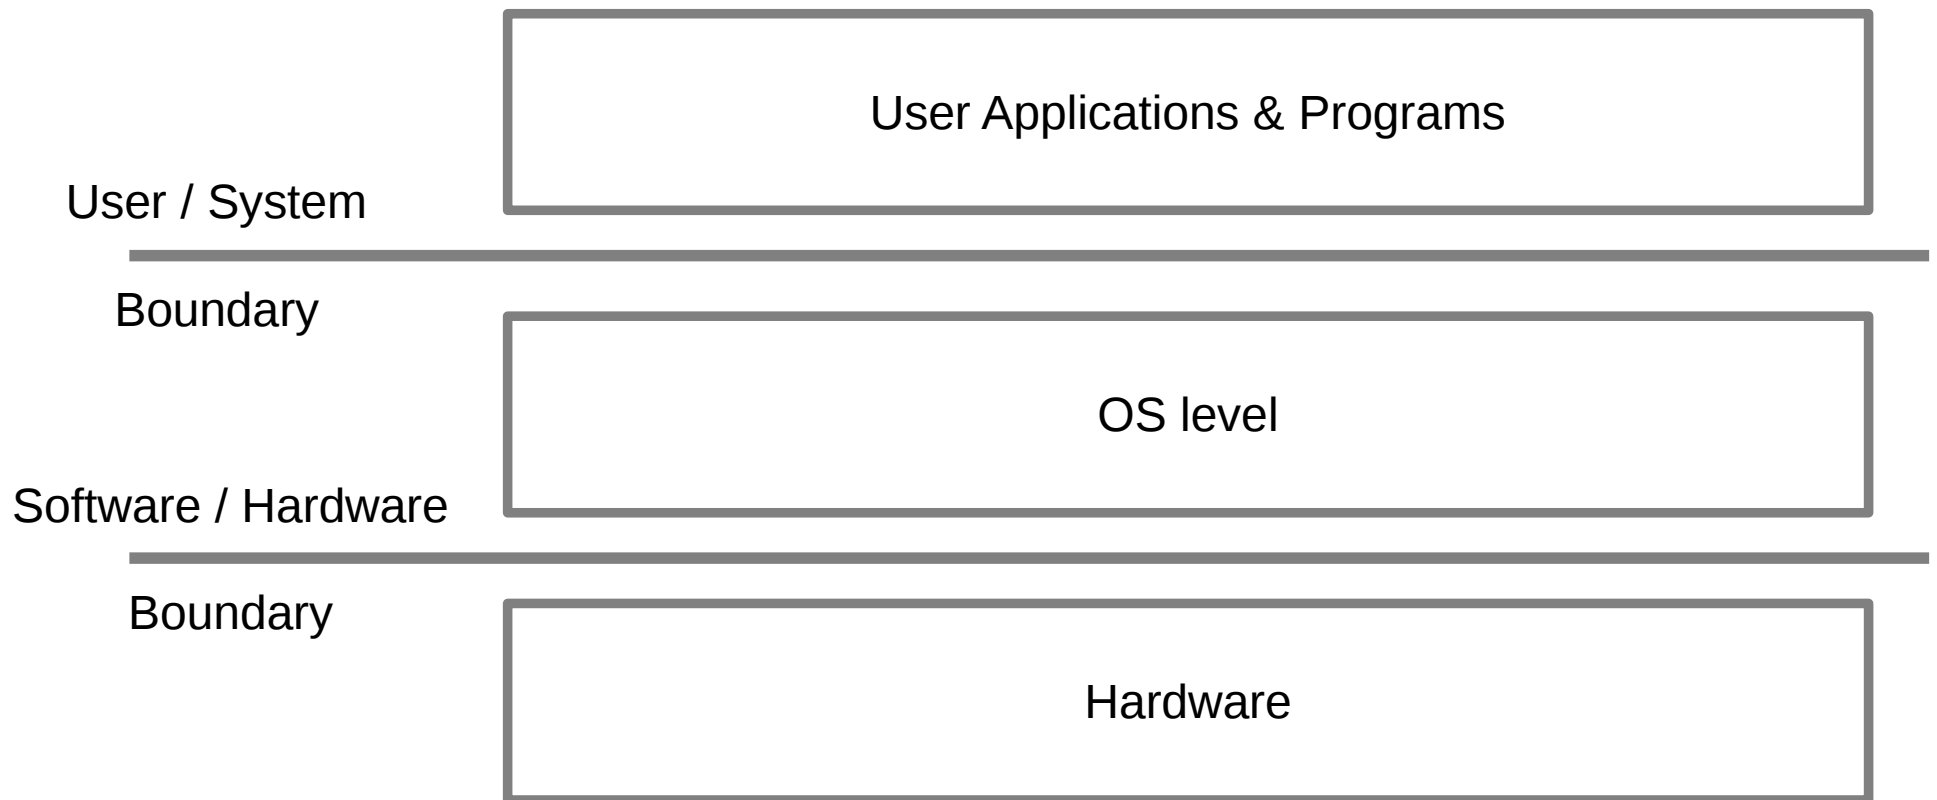
Dimitrios Stamoulis

TR 4110

September 22, 2014

ECSE 420 Tutorials – Introduction

- Why we are here??



ECSE 420 Tutorials – Introduction

- Why we are here??

Different
programming tools

User Applications & Programs

Parallel
programming models

OS level

Software / Hardware

Boundary

Hardware

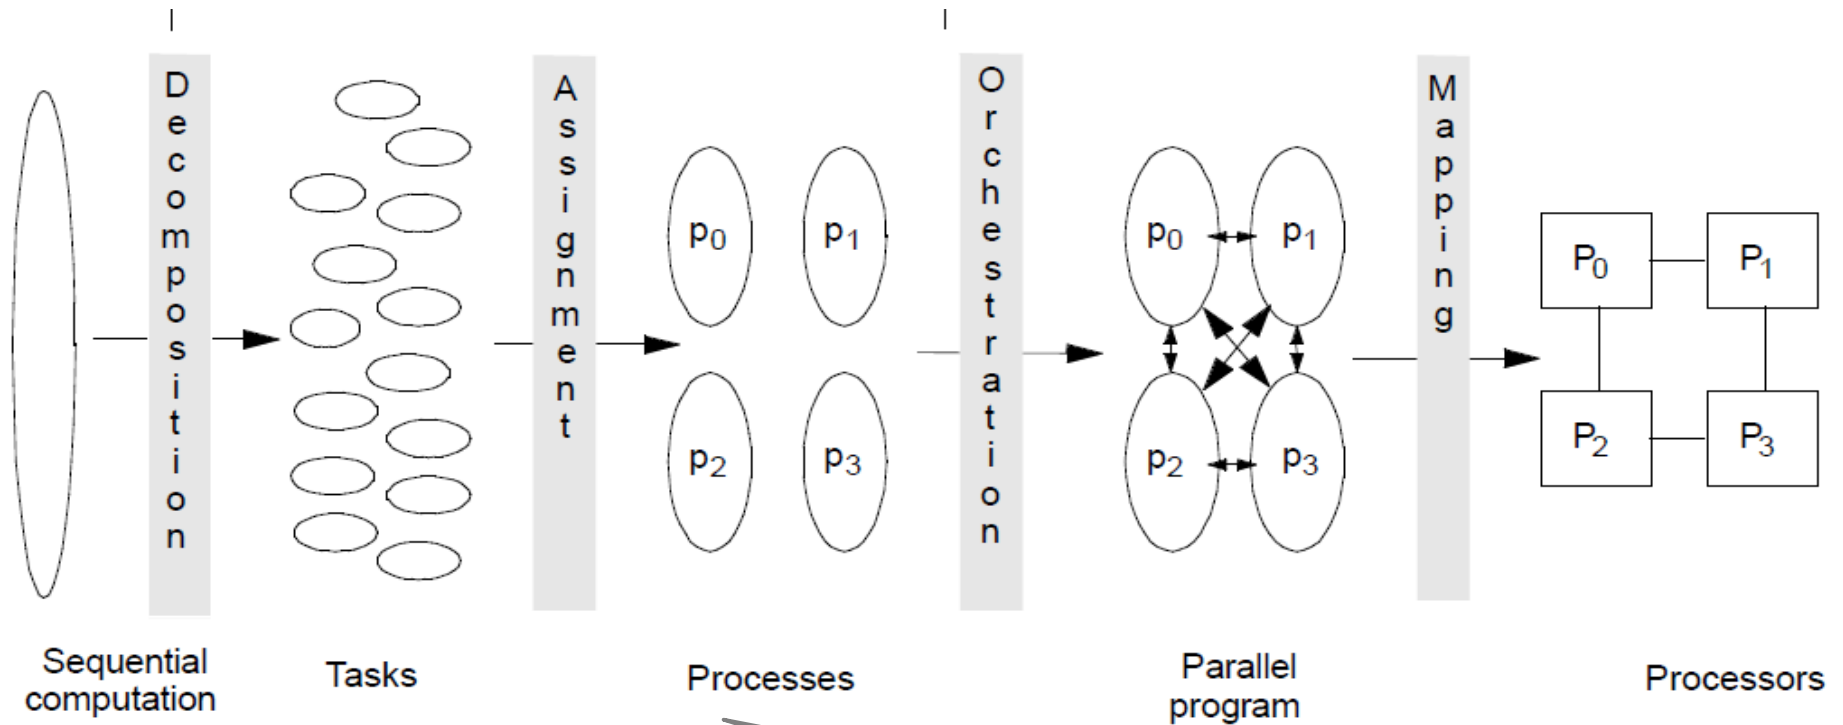
ECSE 420 – Tutorials (1/2)

- Hours and Location
 - Group 1 : Mondays, 04:00 PM - 5:30 PM (TR4110)
 - Group 2 : Wednesdays, 04:00 PM - 5:30 PM (TR4110)
- Our goal to have a foretaste of:
 - Useful parallel programming tools (e.g OpenMP, MPI)
 - Midterm exercises, assignments etc
- So, please ask questions !!
 - Office Hours :
Wednesdays, 10.00-11.30am, McConnell 544
 - Mail them @ :
dimitrios.stamoulis@mail.mcgill.ca

ECSE 420 – Tutorials (2/2)

- Labs Schedule (tentative)
 - Tutorial 1 - 09/22 :
Processes & Threads (Assignment 1)
 - Tutorial 2 - 09/29 :
Processes & Threads (Lab 1)
 - Tutorial 3 - 10/06 :
Shared Memory/ Msg Passing (Assignment 2)
 - Tutorial 4 - 10/13 :
Shared Memory/ Msg Passing (Lab 2)
 - Tutorial 5 - 10/20 :
Parallel/ Distributed Program. (Assignment 3)
 - Tutorial 6 – 10/27:
Parallel/ Distributed Program. (Lab 3)

Tutorial 1 - Introduction



What's a process??

Process

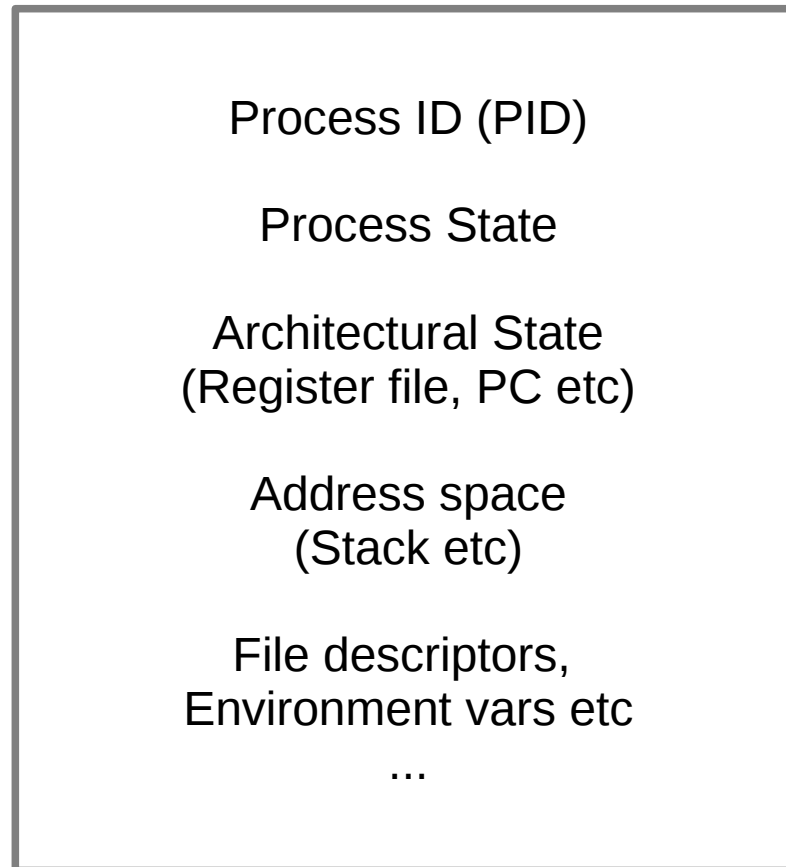
- An 'alive' program
 - currently executed
 - inside an allocated memory portion (RAM)



- Process = program (code) + state
Many processes can execute the same code, but they have different execution states !!
- Unique Process ID (PID)

Process' State

- Process Control Block (PCB)

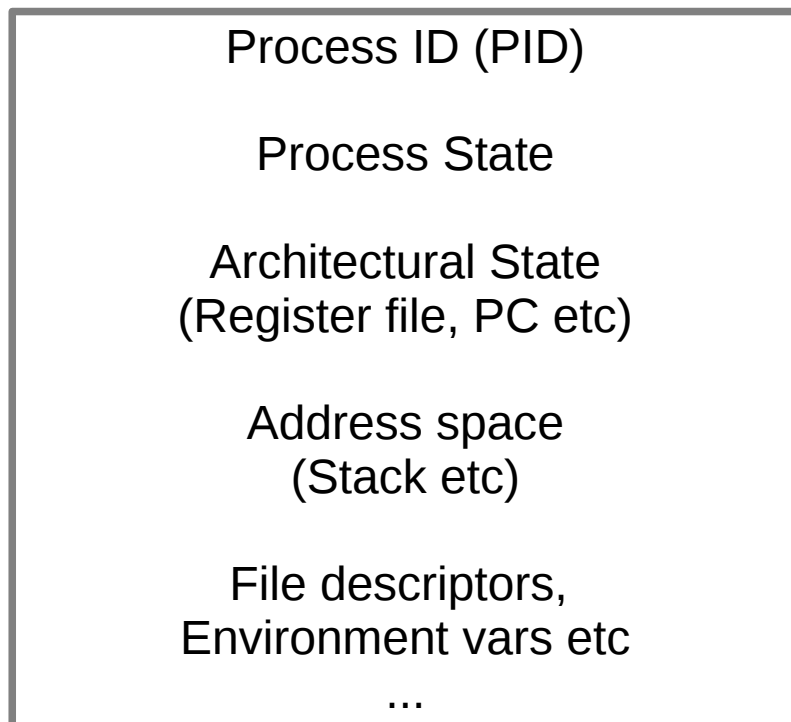


How to create a process : fork()

PID=832 ~ Parent

- Program

- State



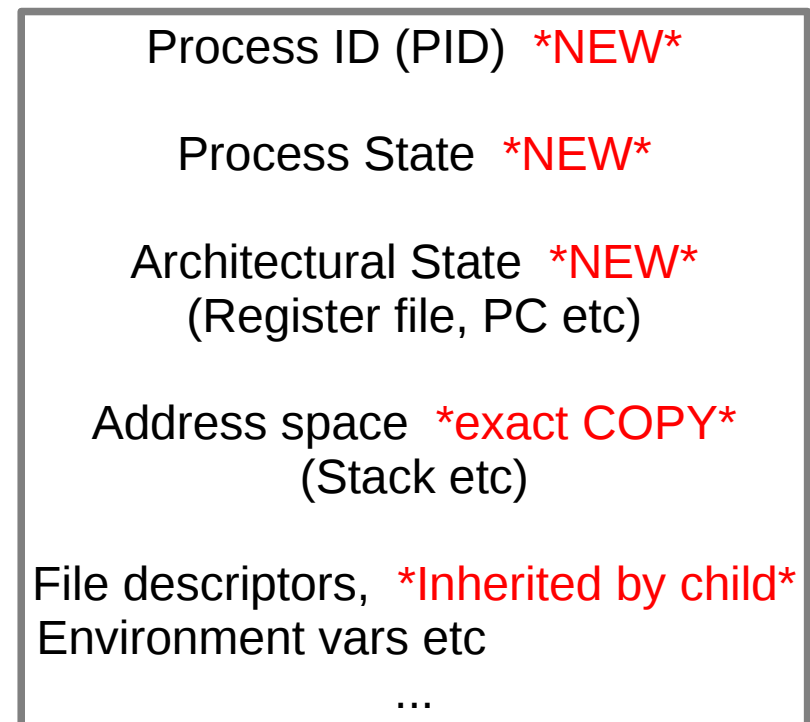
Same program



PID=864 ~ Child

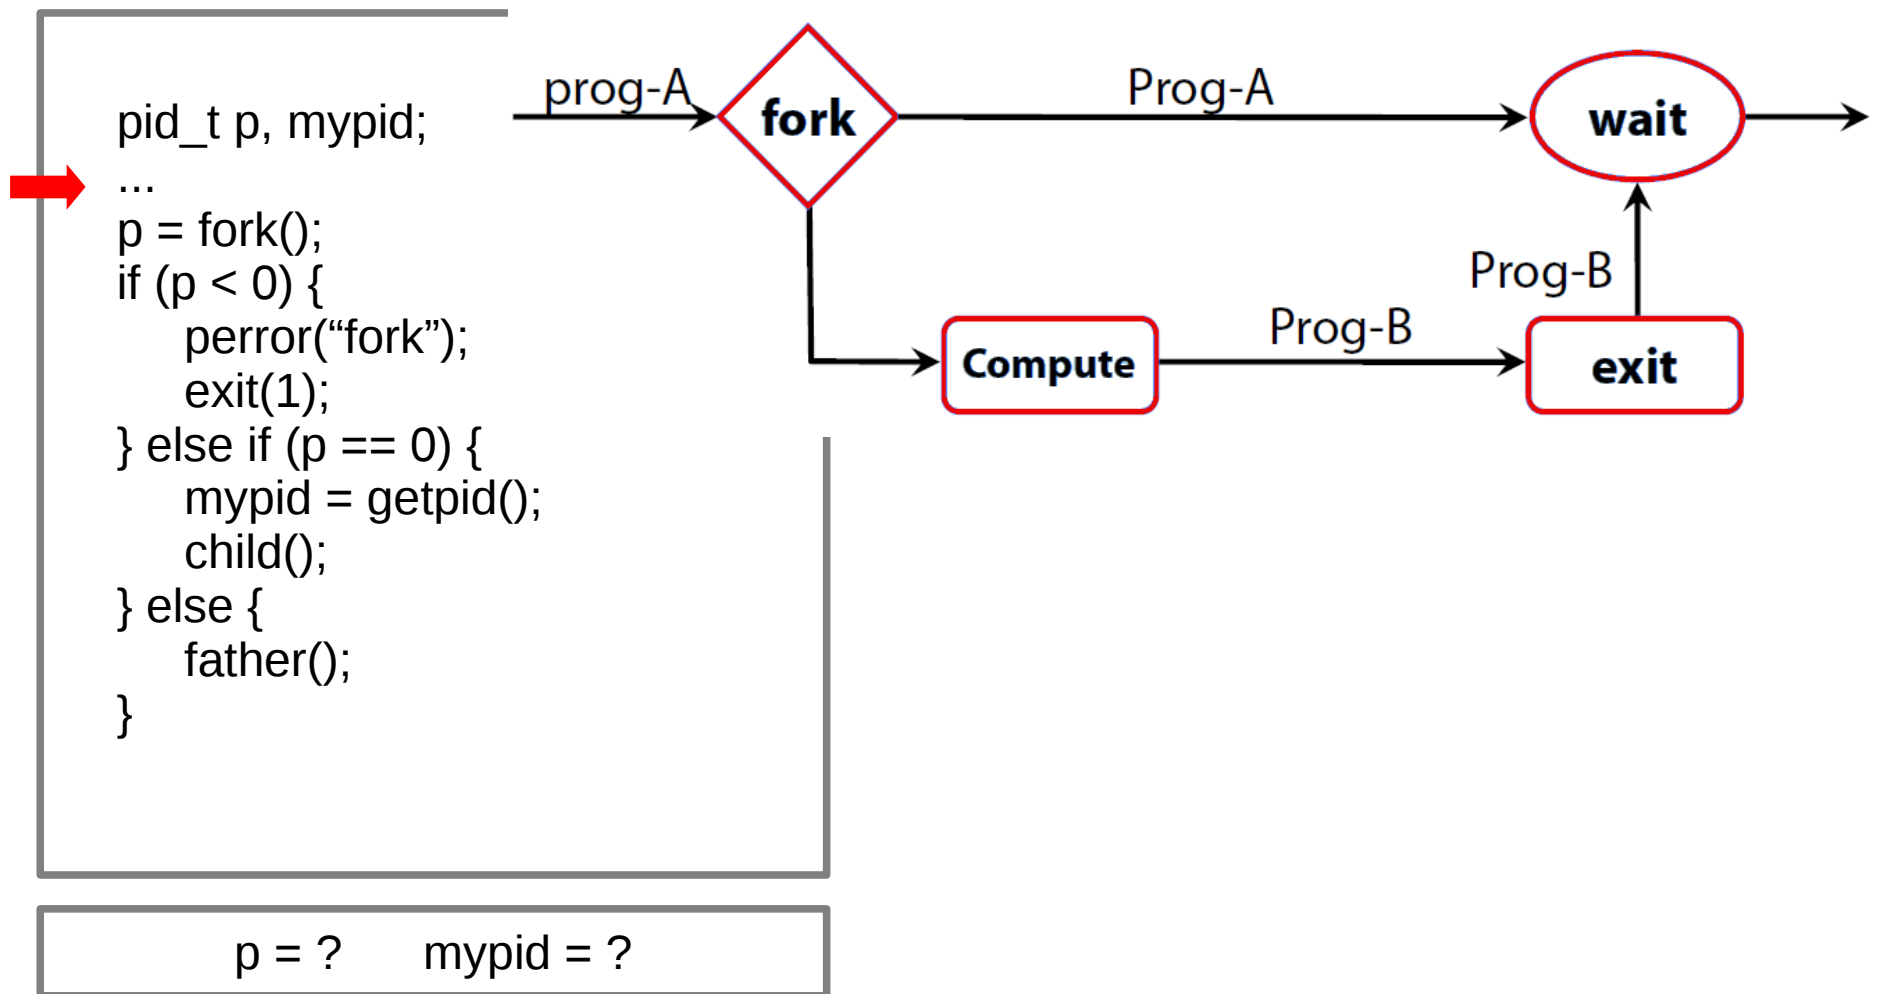
- Program

- State




How to create a process : fork()

PID=832



How to create a process : fork()

PID=832



```
pid_t p, mypid;  
...  
p = fork();  
if (p < 0) {  
    perror("fork");  
    exit(1);  
} else if (p == 0) {  
    mypid = getpid();  
    child();  
} else {  
    father();  
}
```

p = ? mypid = ?

How to create a process : fork()

PID=832


```
pid_t p, mypid;  
...  
p = fork();  
if (p < 0) {  
→ perror("fork");  
  exit(1);  
} else if (p == 0) {  
  mypid = getpid();  
  child();  
} else {  
  father();  
}
```

WHY??

p = -1 mypid = ?

How to create a process : fork()

PID=832




```
pid_t p, mypid;  
...  
p = fork();  
if (p < 0) {  
    perror("fork");  
    exit(1);  
} else if (p == 0) {  
    mypid = getpid();  
    child();  
} else {  
    father();  
}
```

p = ? mypid = ?

How to create a process : fork()

PID=832



```
pid_t p, mypid;  
...  
p = fork();  
if (p < 0) {  
    perror("fork");  
    exit(1);  
} else if (p == 0) {  
    mypid = getpid();  
    child();  
} else {  
    father();  
}
```

p = ? mypid = ?

How to create a process : fork()

PID=832

```
pid_t p, mypid;
...
→ p = fork();
if (p < 0) {
    perror("fork");
    exit(1);
} else if (p == 0) {
    mypid = getpid();
    child();
} else {
    father();
}
```

p = 864 mypid = ?


PID=864

```
pid_t p, mypid;
...
→ p = fork();
if (p < 0) {
    perror("fork");
    exit(1);
} else if (p == 0) {
    mypid = getpid();
    child();
} else {
    father();
}
```

p = 0 mypid = ?

How to create a process : fork()


PID=832



```
pid_t p, mypid;
...
p = fork();
if (p < 0) {
    perror("fork");
    exit(1);
} else if (p == 0) {
    mypid = getpid();
    child();
} else {
    father();
}
```

p = 864 mypid = ?

PID=864




```
pid_t p, mypid;
...
p = fork();
if (p < 0) {
    perror("fork");
    exit(1);
} else if (p == 0) {
    mypid = getpid();
    child();
} else {
    father();
}
```

p = 0 mypid = ?


How to create a process : fork()

PID=832

```
pid_t p, mypid;
...
p = fork();
if (p < 0) {
    perror("fork");
    exit(1);
} else if (p == 0) {
    mypid = getpid();
    child();
} else {
     father();
}
```

p = 864 mypid = ?


PID=864

```
pid_t p, mypid;
...
 p = fork();
if (p < 0) {
    perror("fork");
    exit(1);
} else if (p == 0) {
    mypid = getpid();
    child();
} else {
    father();
}
```

p = 0 mypid = ?


How to create a process : fork()

PID=832

```
pid_t p, mypid;
...
p = fork();
if (p < 0) {
    perror("fork");
    exit(1);
} else if (p == 0) {
    mypid = getpid();
    child();
} else {
     father();
}
```

p = 864 mypid = ?


PID=864

```
pid_t p, mypid;
...
p = fork();
 if (p < 0) {
    perror("fork");
    exit(1);
} else if (p == 0) {
    mypid = getpid();
    child();
} else {
    father();
}
```

p = 0 mypid = ?


How to create a process : fork()

PID=832

```
pid_t p, mypid;
...
p = fork();
if (p < 0) {
    perror("fork");
    exit(1);
} else if (p == 0) {
    mypid = getpid();
    child();
} else {
     father();
}
```

p = 864 mypid = ?


PID=864

```
pid_t p, mypid;
...
p = fork();
if (p < 0) {
    perror("fork");
    exit(1);
} else if (p == 0) {
     mypid = getpid();
    child();
} else {
    father();
}
```

p = 0 mypid = 864


How to create a process : fork()

PID=832

```
pid_t p, mypid;
...
p = fork();
if (p < 0) {
    perror("fork");
    exit(1);
} else if (p == 0) {
    mypid = getpid();
    child();
} else {
     father();
}
```

p = 864 mypid = ?

PID=864

```
pid_t p, mypid;
...
p = fork();
if (p < 0) {
    perror("fork");
    exit(1);
} else if (p == 0) {
    mypid = getpid();
     child();
} else {
    father();
}
```

p = 0 mypid = 864

Example - fork()


```
int main()
{
    pid_t fork(void);
    pid_t pid, child_pid, parent_pid;

    switch (pid = fork())
    {
        case -1:
            perror("The fork failed!");
            break;

        case 0:
            child_pid = getpid();
            printf("\nHello I am a new Child w/ PID %d\n", child_pid);
            break;

        default:
            parent_pid = getpid();
            printf("\nHello I am parent w/ PID %d and\nI created the process %d\n", parent_pid, pid);
    }

    return 0;
}
```




How to create a process : fork()

PID=832 ~ Parent

- Program

- State




Same program



PID=864 ~ Child

- Program

- State



Example - fork_mem()

```
int main()
{
    pid_t pid, child_pid;
    int a[10]= {0,10,20,30,40,50,60,70,80,90};
    int i;


    switch (pid = fork())
    {
        case -1:
            perror("The fork failed!");
            break;

        case 0:
            child_pid = getpid();
            printf("Hello I am a new Child my pid is %d\n",child_pid);
            a[4]= 50;
            break;

        default:
            printf("Hello I am parent I just created a new process %d\n",pid);
            a[4]= 30;
    }

    for ( i = 0; i < 10; i++)
        printf("a[%d]=%d\n", i , a[i]);


    return 0;
}
```




Inter-Process Communication (IPC)

IPC Mechanisms


- Signals




- Pipes



- POSIX messages



- Shared Memory Segments



Example - wait()

- Execute `$./wait`

Parent process started at Mon Sep 30 02:47:07 2013


Parent waiting for child at Mon Sep 30 02:47:07 2013

Child w/ PID 8580 started at Mon Sep 30 02:47:07 2013

Parent waiting for child at Mon Sep 30 02:47:08 2013

...

Child ended normally



- What if i stop the child?

- `kill -STOP 8580`

- What if i kill the child?

- `kill -9 8580`


- What if i kill the parent?

I can use a signal as an execution checkpoint...!!

We can extend this idea to interprocess messages :

`int MPI_Bcast()`

Shared memory → Critical Sections



```
15.      while (!done) do                                /*outer loop sweeps*/
16.          mydiff = diff = 0;                          /*set global diff to 0 (okay for all to do it)*/
16a.      BARRIER(bar1, nprocs);                       /*ensure all reach here before anyone modifies diff*/
17.          for i ← mymin to mymax do                   /*for each of my rows*/
18.              for j ← 1 to n do                       /*for all nonborder elements in that row*/
19.                  temp = A[i,j];
20.                  A[i,j] = 0.2 * (A[i,j] + A[i,j-1] + A[i-1,j] +
21.                                A[i,j+1] + A[i+1,j]);
22.                  mydiff += abs(A[i,j] - temp);
23.              endfor
24.          endfor
25a.      LOCK(diff_lock);
25b.          diff += mydiff;
25c.      UNLOCK(diff_lock);
25d.      BARRIER(bar1, nprocs);
25e.      if (diff/(n*n) < TOL) then done
25f.      BARRIER(bar1, nprocs);
```


A red arrow points from line 25a to the `LOCK(diff_lock);` statement.

We will see same techniques :

- OpenMP :
`#pragma omp critical`

Threads (1/2)

- A thread exists within a process :
A finer-grained unit of execution than processes
- \$./myprogram → Linux creates a new process →
Linux creates a single, master thread.



- The master thread can create additional threads.

Threads (2/2)

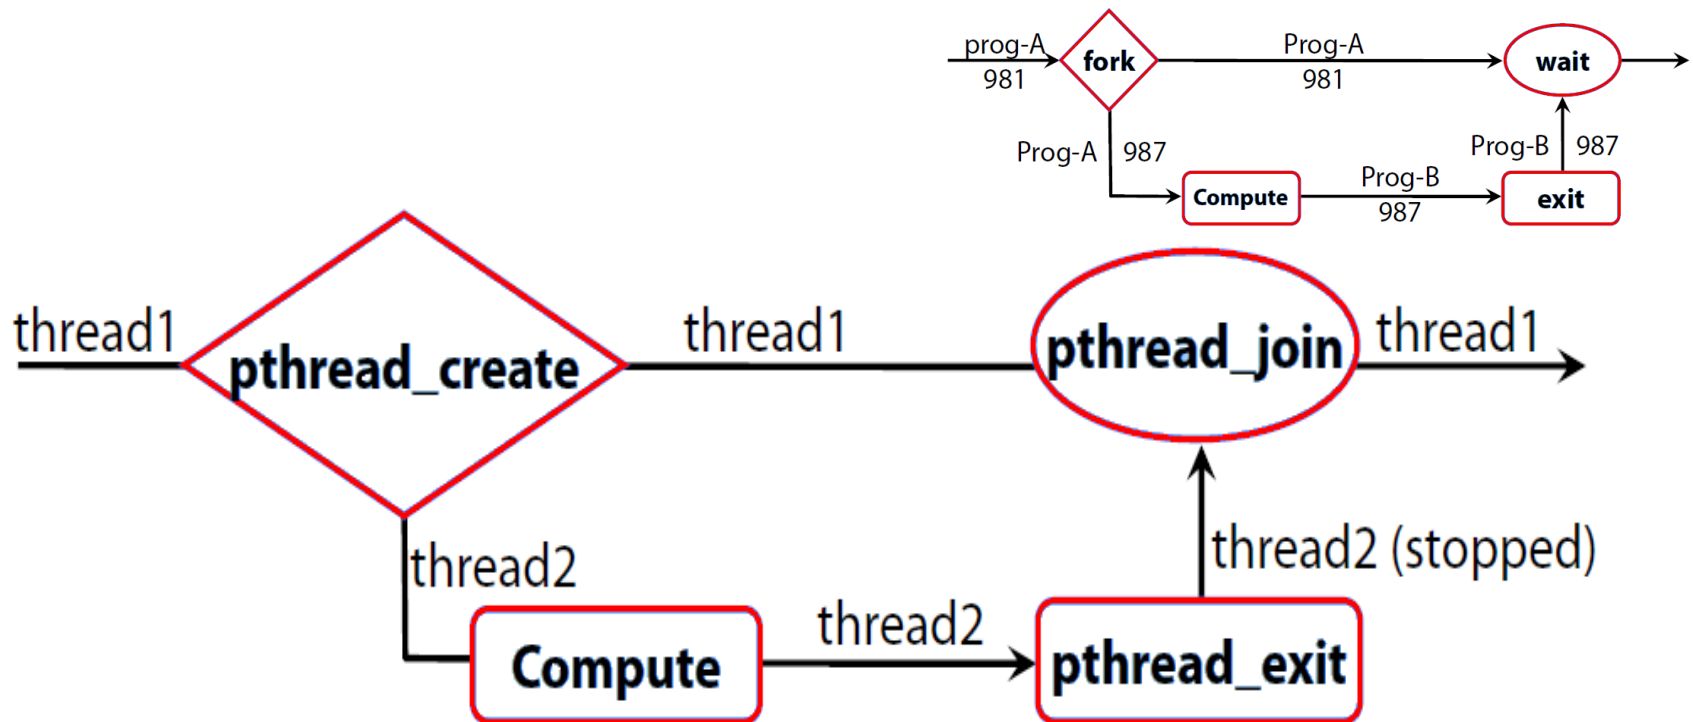
- Threads within a process use process' resources :
 - Share process' resources
Is it always good ? → NO !!
WHY? → **explicit synchronization by the programmer**
 - Duplicate only architectural state
WHY? → To be independently schedulable

Threads vs Processes

- Thread creation is “lightweight” - WHY?
 - A thread uses less resources than a process.
- Less communication overhead
- Less scheduling overhead
- More efficient programming techniques / tools :
 - e.g : OpenMP

```
#pragma omp parallel num_threads(4)
{
    ....
}
```

Pthreads (POSIX threads) - pthread_create()



`pthread_create(thread, attr, start_routine, arg)` arguments:

- **thread**: Identifier for the new thread.
- **Attr**: Attribute object (e.g Stack size, priority). NULL for the default values!!
- **start_routine**: the C routine that the thread will execute.
- **arg**: A single argument that may be passed to **start_routine**.

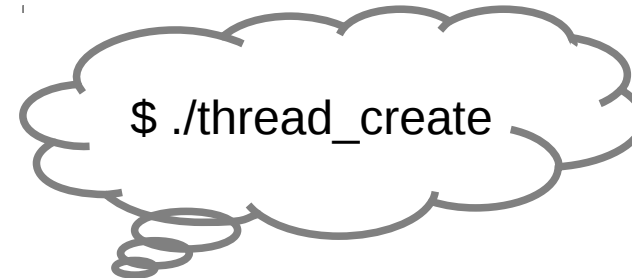
How to create a thread : pthread_create()

```
#include <pthread.h>
#include <stdio.h>
#include <stdlib.h>
#define NUM_THREADS 5

void *PrintHello(void *threadid)
{
    long tid;
    tid = (long)threadid;
    printf("Hello World! It's me, thread #%ld!\n", tid);
    pthread_exit(NULL);
}

int main (int argc, char *argv[])
{
    pthread_t threads[NUM_THREADS];
    int rc;
    long t;
    for(t=0; t<NUM_THREADS; t++){
        printf("In main: creating thread %ld\n", t);
        rc = pthread_create(&threads[t], NULL, PrintHello, (void *)t);
        if (rc){
            printf("ERROR; return code from pthread_create() is %d\n", rc);
            exit(-1);
        }
    }

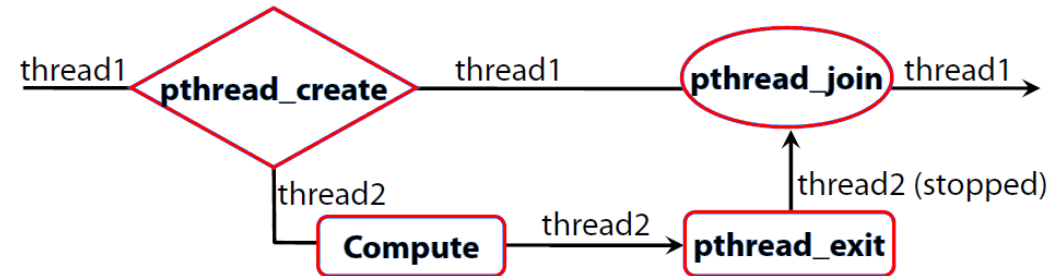
    /* Last thing that main() should do */
    pthread_exit(NULL);
}
```



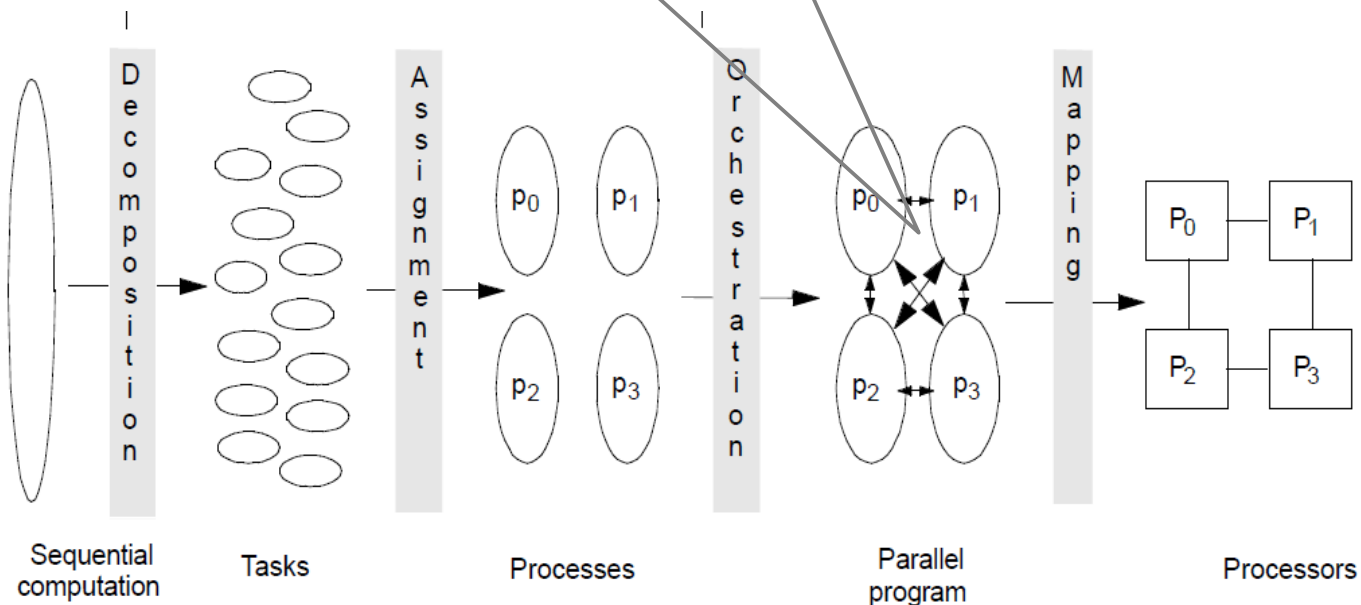
Pthreads (POSIX threads) - pthread_join()

pthread_join(threadid,status) :

- Gives the **threadid** termination return status if it was specified in the **threadid** call to pthread_exit().
- pthread_join() blocks the calling thread until the specified **threadid** terminates.

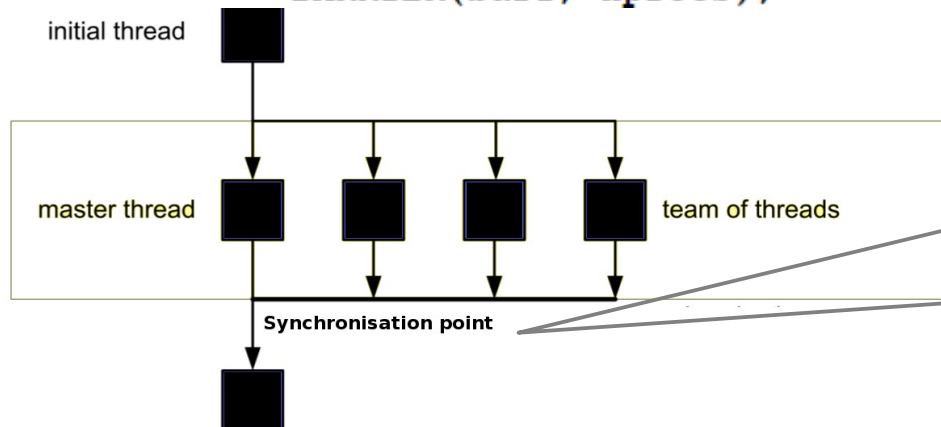


WHY is that important? → **Synchronization Point !!**



pthread_join() - Synchronization

```
15.      while (!done) do                                /*outer loop sweeps*/
16.          mydiff = diff = 0;                          /*set global diff to 0 (okay for all to do it)*/
16a.      BARRIER(bar1, nprocs);                       /*ensure all reach here before anyone modifies diff*/
17.          for i ← mymin to mymax do                   /*for each of my rows*/
18.              for j ← 1 to n do                       /*for all nonborder elements in that row*/
19.                  temp = A[i,j];
20.                  A[i,j] = 0.2 * (A[i,j] + A[i,j-1] + A[i-1,j] +
21.                      A[i,j+1] + A[i+1,j]);
22.                  mydiff += abs(A[i,j] - temp);
23.              endfor
24.          endfor
25a.      LOCK(diff_lock);                               /*update global diff if necessary*/
25b.          diff += mydiff;
25c.      UNLOCK(diff_lock);
25d.      BARRIER(bar1, nprocs);                       /*ensure all reach here before checking if done*/
25e.      if (diff/(n*n) < TOL) then done = 1; /*check convergence; all get
                                                    same answer*/
25f.      BARRIER(bar1, nprocs);
```



We will see same techniques :

- OpenMP :
 [#pragma omp barrier](#)
- MPI :
 [int MPI_Barrier\(\)](#)

Joining threads : pthread_join()

```
int main (int argc, char *argv[])
{
    pthread_t thread[NUM_THREADS];
    int rc;
    long t;
    void *status;

    for(t=0; t<NUM_THREADS; t++) {
        printf("Main: creating thread %ld\n", t);
        rc = pthread_create(&thread[t], NULL, BusyWork, (void *)t);
        if (rc) {
            printf("ERROR; return code from pthread_create() is %d\n", rc);
            exit(-1);
        }
    }

    for(t=0; t<NUM_THREADS; t++) {
        rc = pthread_join(thread[t], &status);
        if (rc) {
            printf("ERROR; return code from pthread_join() is %d\n", rc);
            exit(-1);
        }
        printf("Main: completed join with thread %ld having a status of %ld\n", t, (long)status);
    }

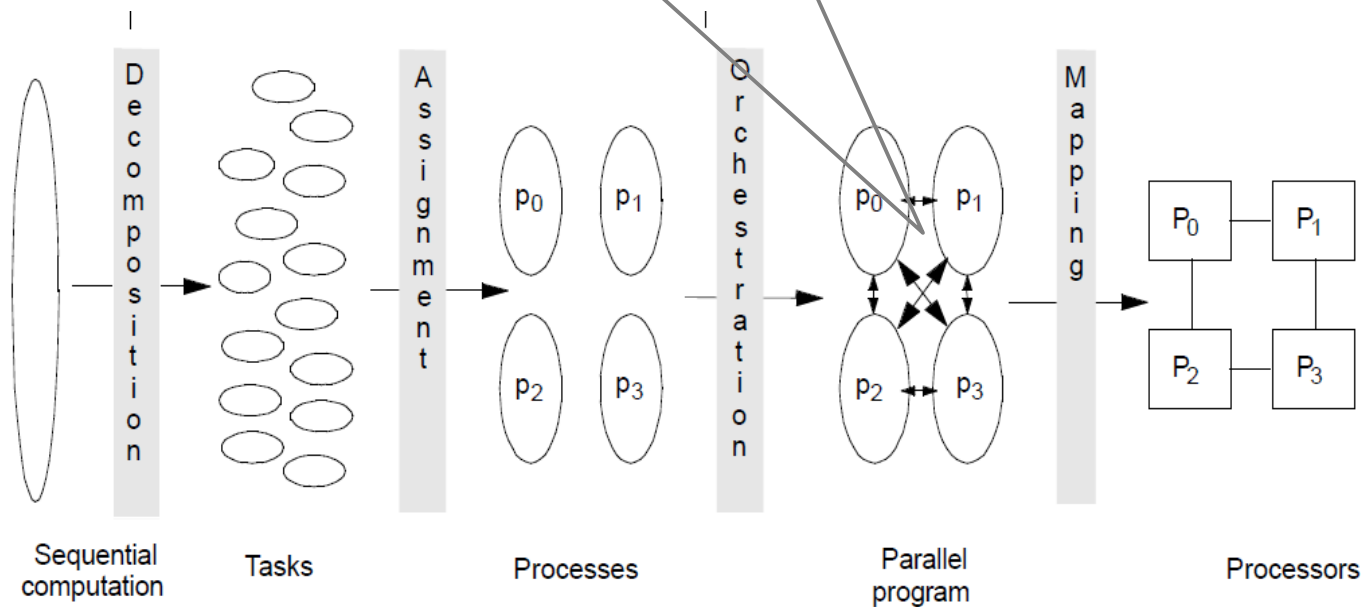
    printf("Main: program completed. Exiting.\n");
    pthread_exit(NULL);
}
```



Synchronization (1/2)

WHY is that important? →


Synchronization Point !!




Synchronization (2/2)

Inter-Process Communication (IPC) Mechanisms


- Signals



- Pipes



- POSIX messages



- Shared Memory Segments

