

OpenMP


ECSE 420 - Tutorial 3

Dimitrios Stamoulis

TR 4110

October 6, 2014

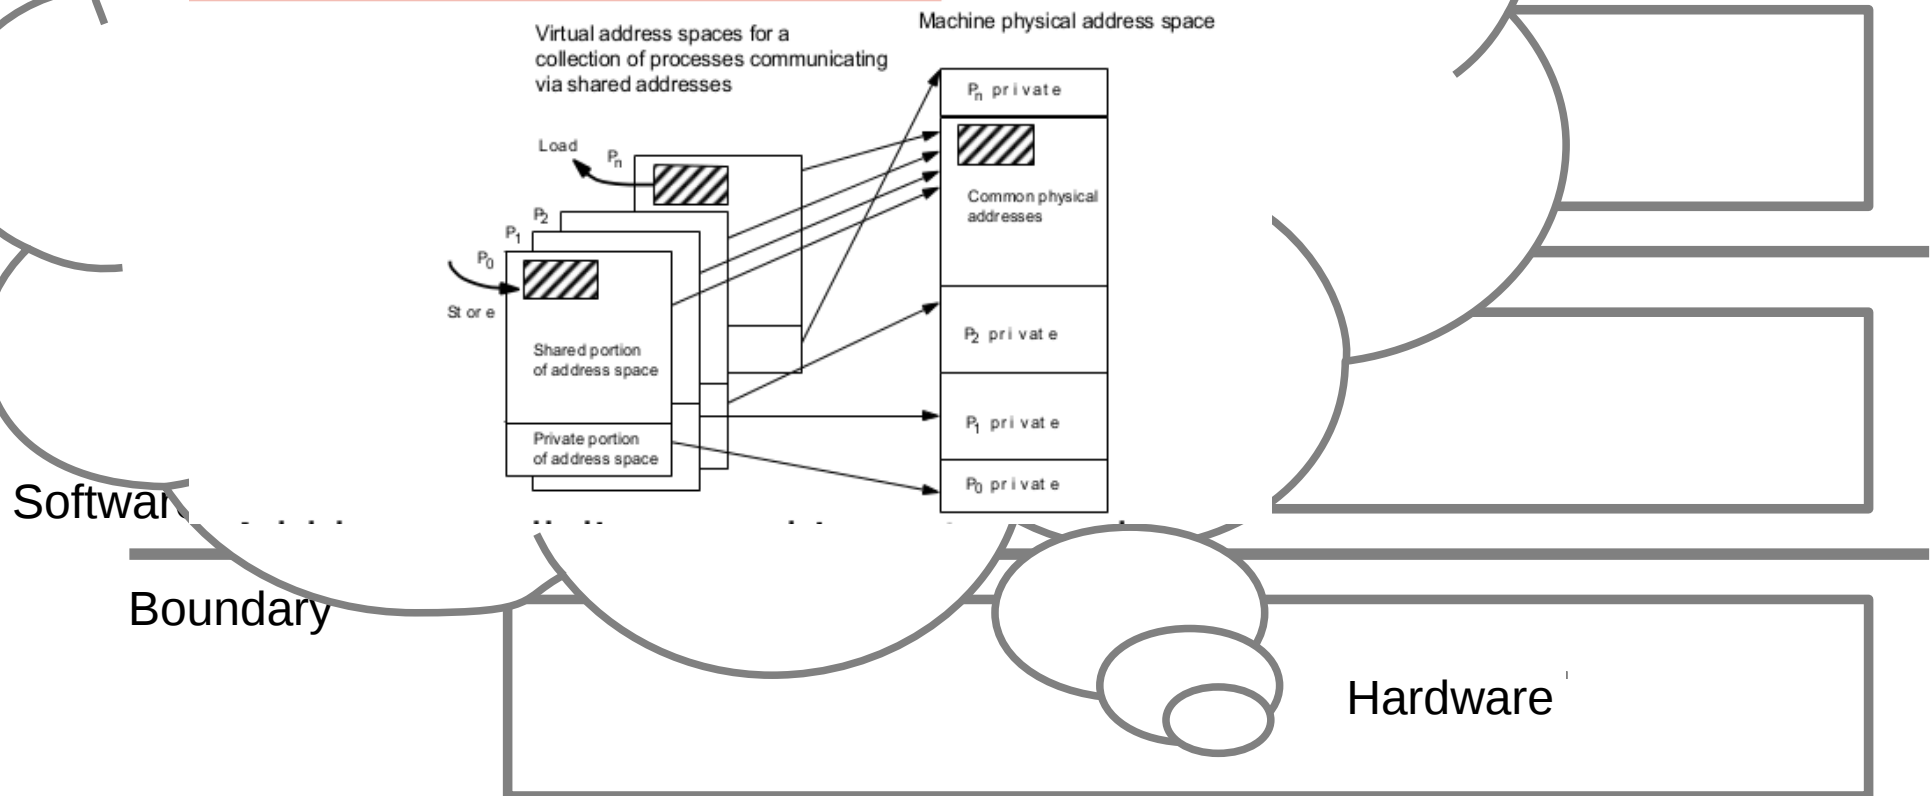
Tutorial 3 – Introduction



Tutorial 3 – Introduction

Structured Shared Address Space

From ECSE 420
lecture slides!!



Tutorial 3 – Introduction

- Why we are here??

OpenMP

User Applications & Programs

Shared memory
multiprocessing
programming tools

OS level

Software / Hardware

Boundary

Hardware

Tutorial 3 - Material

- Linux machine :
 - Compile an OpenMP program :
\$ gcc **-fopenmp** myprogram.c
\$./a.out

Shared Variable - A simple example

```
int main()
{
    int sharedVar = 5;
    CREATE THREAD( thread1 )
    CREATE THREAD( thread2 )
    WAIT THREADS(..)
    printf ("shared var = %d\n", sharedVar);
    return 0;
}
```

Thread 1:

```
for (i=0; i<LOOPS; i++)
    *shared_var *= 2;
```

Thread 2:

```
for (i=0; i<LOOPS; i++)
    *shared_var -= 2;
```



```
stam@Stam:~/codes$ for i in $(seq 5); do ./sharedVar ; done
shared variable=-36700160
shared variable=5242840
shared variable=5242840
shared variable=5242840
shared variable=-36700160
```

OpenMP (Open Multi-Processing)

- Application programming interface (API)
- Shared memory multiprocessing programming API:
 - Run-time library
 - Compiler directives (**pragmas**)
 - Environment variables

API → Sets the behaviors of the **runtime system**
→ The runtime system implements the **execution model**

OpenMP (Open Multi-Processing)

- Application programming interface (API)
- Shared memory multiprocessing programming API:
 - Run-time library
 - Compiler directives (**pragmas**)
 - Environment variables

API → Sets the behaviors of the **runtime system**

→ The runtime system implements the

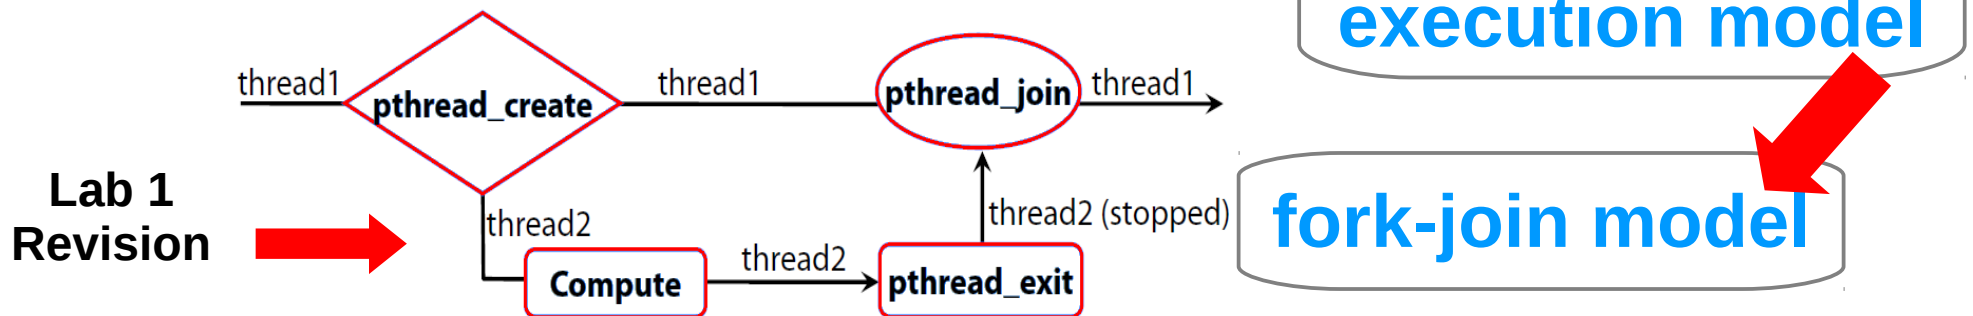
OpenMP
execution model

OpenMP (Open Multi-Processing)

- Application programming interface (API)
- Shared memory multiprocessing programming API:
 - Run-time library
 - Compiler directives (**pragmas**)
 - Environment variables

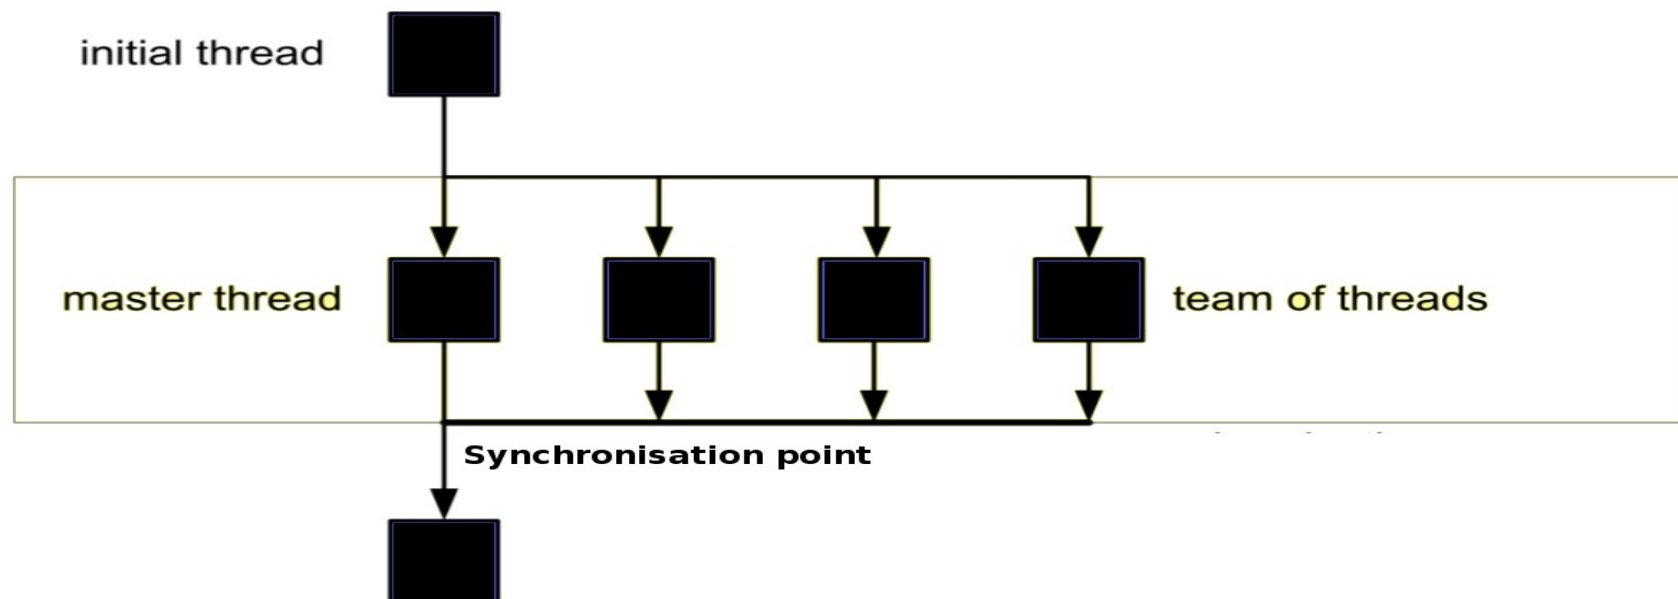
API → Sets the behaviors of the **runtime system**

→ The runtime system implements the




Execution model (1/2)

- **OpenMP uses a fork-join model of parallel execution :**
 - A thread encounters an '**omp parallel**' section
 - The thread creates a team composed of itself and other additional threads
 - The encountering thread becomes the master of the new team
 - All team members execute the code inside the '**omp parallel**' section



Execution model (1/2)



- **OpenMP uses a fork-join model of parallel execution.**
 - All threads finish their work and wait at the implicit barrier at the end of the '**omp parallel**' section.
 - When **ALL** team members have arrived at the barrier, the threads can leave the barrier.
 - The master thread continues execution of user code beyond the end of the parallel section.

Basic Concepts

- OpenMP API programming tools:

- Run-time library
- Compiler directives
- Environment variables

- Number of threads
- Thread ID
- Timers
- ...

- Number of threads
- Scheduling type
- Nested parallelism
- ...

- Parallel regions
- Work sharing
- Synchronization
- Specific attributes
 - private
 - shared
 - Reduction
- ...

Basic Concepts

- I can do the same thing with more than one way :

- Run-time library
- Compiler directives
- Environment variables



```
omp set_num_threads(4)
```

```
#pragma omp parallel num_threads(4)
```

```
export OMP_NUM_THREADS=4
```

Basic Concepts – Compiler directives

#pragma omp <directive> <clauses>

Different <directive> types to:

- Define Parallel region
 - #pragma omp parallel
- Orchestrate Work sharing
 - #pragma omp for
 - #pragma omp single
- Achieve Synchronization
 - #pragma omp barrier
 - #pragma omp master
 - #pragma omp critical
 - #pragma omp ordered

Basic Concepts – Compiler directives

#pragma omp <directive> <clauses>

Different <directive> types to:

- Define Parallel region
 - #pragma omp parallel
- Orchestrate Work sharing
 - #pragma omp for
 - #pragma omp single
- Achieve Synchronization
 - #pragma omp barrier
 - #pragma omp master
 - #pragma omp critical
 - #pragma omp ordered



Parallel region

```
#pragma omp parallel <clauses>  
{ structured-block }
```

We can define different <clauses> :



- num_threads
- private variables
- shared variables
- ...

We have an implicit barrier at the end of the 'omp parallel' section → The barrier implies 'flush'

Parallel region - Our 1st example ...

```
#include <omp.h>

int main(){
  ...
  #pragma omp parallel num_threads(4)
  {
    ... // our parallel section
  }
  ...
}
```



Parallel region - Our 2nd example ...

```
#include <omp.h>
```

```
int main(){
```

```
...
```

```
#pragma omp parallel num_threads(4)
```

```
{  
... // our parallel section #1  
}
```

```
...
```

```
omp_set_num_threads(3);
```

```
#pragma omp parallel
```

```
{  
... // our parallel section #2  
}
```

```
...
```

```
}
```

initial thread

master thread

team of threads


Synchronisation point

initial thread

master thread

team of threads

Synchronisation point

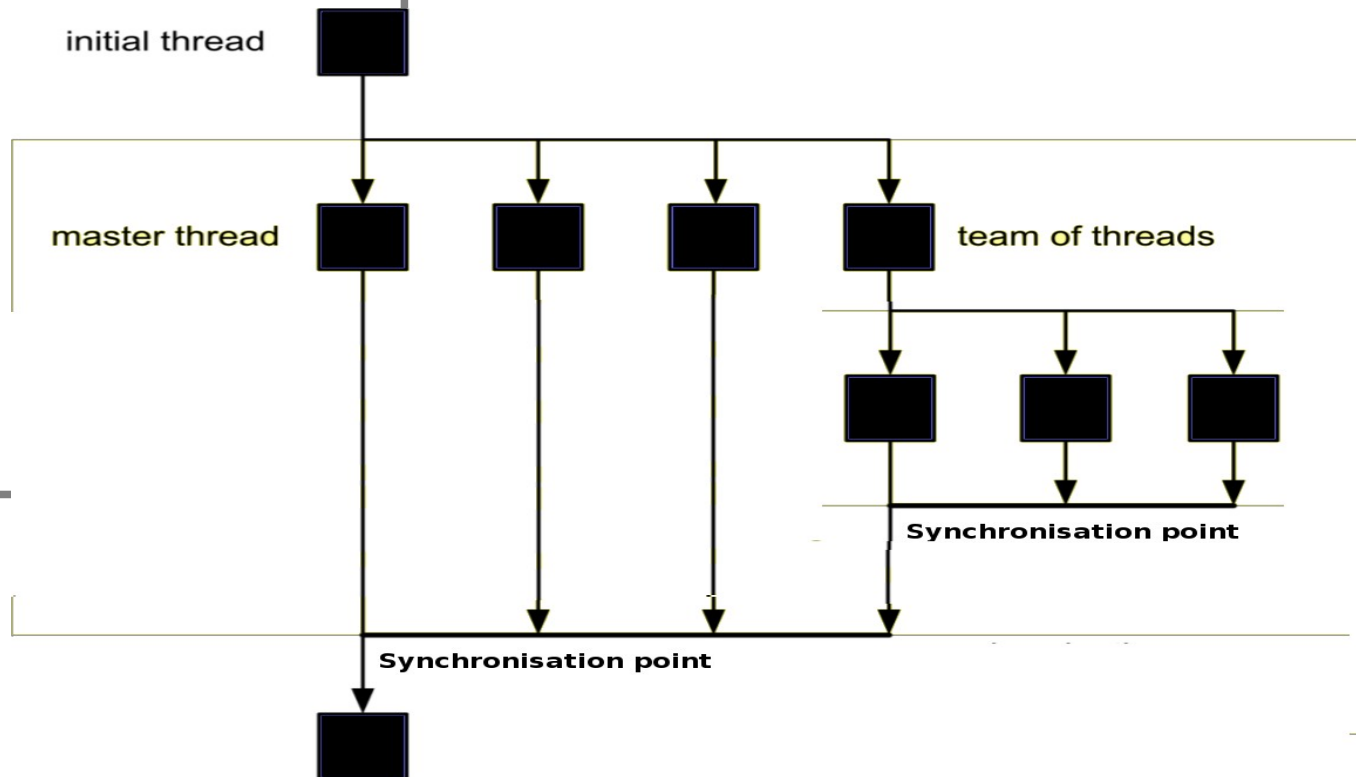


Nested parallel regions (1/2)

```
#include <omp.h>

int main(){
    ...
    omp_set_nested(1);
    #pragma omp parallel num_threads(4)
    {
        ... // our parallel section #1

        if (omp_get_thread_num() == 3){
            ...
            omp_set_num_threads(3);
            #pragma omp parallel
            {
                ...
            }
        }
    }
}
```



Nested parallel regions (1/2)

```
#include <omp.h>
```

```
int main(){
```

```
...
```

```
omp_set_nested(1);  
#pragma omp parallel num_threads(4)
```

```
{
```

```
... // our parallel section #1
```

```
if (omp_get_thread_num() == 3){
```

```
...
```

```
omp_set_num_threads(3);  
#pragma omp parallel
```

```
{
```

```
...
```


```
}
```

```
}
```

```
...
```

```
}
```

What if `omp_set_nested(0);` ????



Nested parallel regions (2/2)

```
#include <omp.h>
```

```
int main(){
```

```
...
```

```
omp_set_nested(1);  
#pragma omp parallel num_threads(4)
```

```
{
```

```
... // our parallel section #1
```

```
if (omp_get_thread_num() == 3){
```

```
...
```

```
omp_set_num_threads(3);  
#pragma omp parallel
```

```
{
```

```
...
```

```
}
```

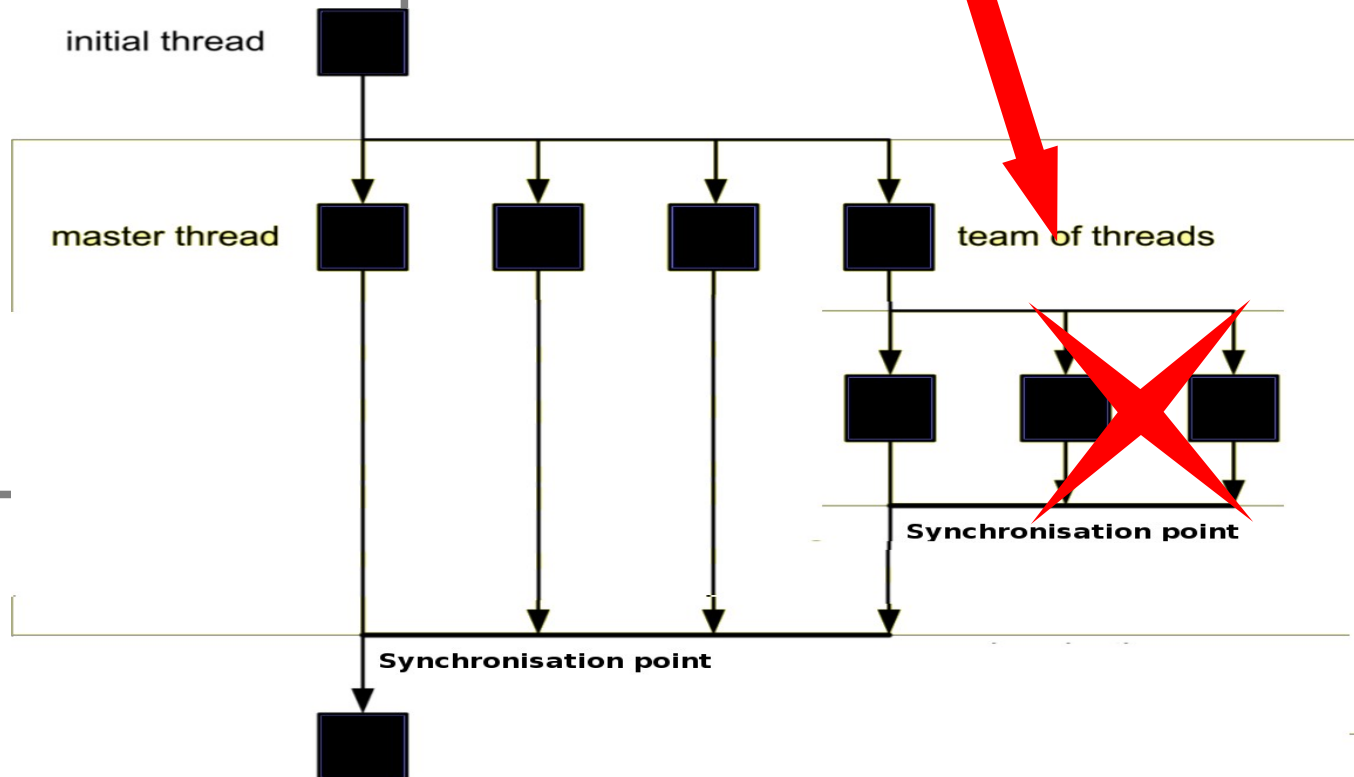
```
}
```

```
...
```

```
}
```

What if `omp_set_nested(0);` ????

setenv OMP_NESTED FALSE



Basic Concepts – Compiler directives

#pragma omp <directive> <clauses>

Different <directive> types to:

- Define Parallel region
 - #pragma omp parallel
- Orchestrate Work sharing
 - #pragma omp for
 - #pragma omp single
- Achieve Synchronization
 - #pragma omp barrier
 - #pragma omp master
 - #pragma omp critical
 - #pragma omp ordered



Work sharing

```
#pragma omp for [schedule(...)] [nowait]  
{ for - loop }
```

We can define different scheduling schemes :

- static : round-robin scheduling
- dynamic : Inactive threads are scheduled
- ...

We can prevent the synchronization at the implicit barrier at the end of the for-loop → **nowait**

Parallel for

```
// Array addition example
int main(){
  ...

  // independent loop iterations
  for (i = 0; i < n; i++)
    c[i] = a[i] + b[i];
  ...
}
```




Which variables shall I share?

Shall I share everything?? (1/2)

```
// Array addition example
int main(){
  ...

  // independent loop iterations
  for (i = 0; i < n; i++)
    c[i] = a[i] + b[i];

  ...
}
```



- Within a parallel section, we can have :
 - Shared variables
 - Private variables !!

Shall I share everything?? (2/2)

```
// independent loop iterations
for (i = 0; i < n; i++)
    c[i] = a[i] + b[i];
```



```
#include <omp.h>

int main(){
    ...
    // independent loop iterations
    #pragma omp parallel shared(n, a, b, c) private(i) num_threads(5)
    #pragma omp for
    for (i = 0; i < n; i++)
        c[i] = a[i] + b[i];
    ...
}
```

Parallel for - nowait

```
#include <omp.h>

int main(){
    ...
    #pragma omp parallel num_threads(2)
    {
        #pragma omp for nowait
        for (i = 0; i < 5; i++)
        {
            sleep(omp_get_thread_num()*2);
        }
        // verify that thread 0 did not wait
        printf("I am thread: %d\n", omp_get_thread_num() );
    }
    ...
}
```

Basic Concepts – Compiler directives

#pragma omp <directive> <clauses>

Different <directive> types to:

- Define Parallel region
 - #pragma omp parallel
- Orchestrate Work sharing
 - #pragma omp for
 - #pragma omp single
- Achieve Synchronization
 - #pragma omp barrier
 - #pragma omp master
 - #pragma omp critical
 - #pragma omp ordered



Synchronization

```
#pragma omp <directive>  
{ structured-block }
```

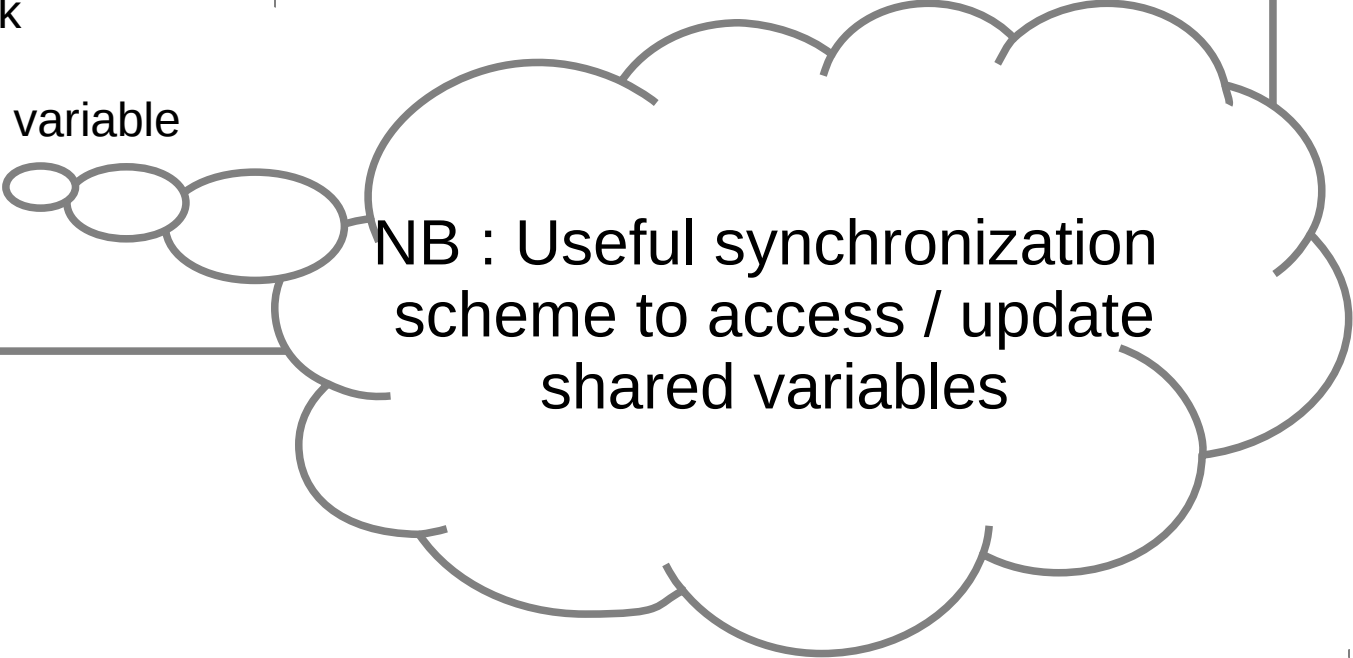
We can define different <directives> :

- #pragma omp critical
- #pragma omp barrier
- #pragma omp master
- #pragma omp flush
- #pragma omp ordered

Example – critical

```
#include <omp.h>

int main(){
    ...
    #pragma omp parallel shared(TotalSum) private(mySum)
    {
        ...
        work(); //let's do our work
        #pragma omp critical
        // let's update our shared variable
        TotalSum += mySum;
    }
}
```



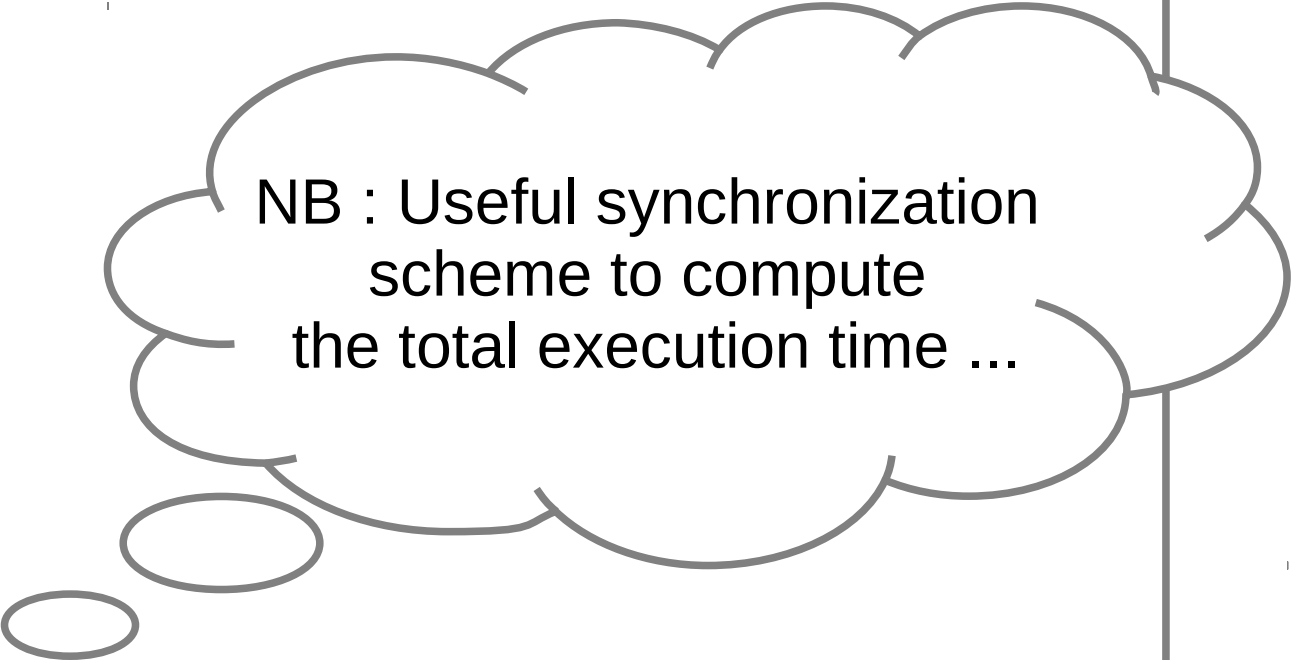
NB : Useful synchronization scheme to access / update shared variables

Example – barrier & master (1/2)

```
#include <omp.h>

int main(){
    ...
    #pragma omp parallel
    {
        #pragma omp barrier
        #pragma omp master
        Gettimeofday(...);

        work(); //let's do our work
        #pragma omp barrier
        #pragma omp master
        {
            gettimeofday(...);
            printf("execution time", ...);
        }
    }
}
```



NB : Useful synchronization scheme to compute the total execution time ...

Example – barrier & master (2/2)


```
#include <omp.h>

int main(){
    ...
    #pragma omp parallel
    {
        #pragma omp barrier
        #pragma omp master
        Gettimeofday(...);


        work(); //let's do our work
        #pragma omp barrier
        #pragma omp master
        {
            gettimeofday(...);
            printf("execution time", ...);
        }
    }
}
```

Suppose we have a program that can be divided into 5 concurrent tasks. One of these tasks requires twice the execution time than the other 4. **Max Speed-up=??**

Do I always share a shared variable !?!?




Do I always share a shared variable !?!?




- Keep in mind that we have machines with complex Memory hierarchy :
 - Multi-level caches.
 - False sharing cases.
 - ...

Do I always share a shared variable !?!?



- Keep in mind that we have machines with complex Memory hierarchy :
 - Multi-level caches.
 - False sharing cases.
 - ...



Example - flush()

```
#include <omp.h>

int main(){
    ...
    // initially a = b = 0 !!
    #pragma omp parallel num_threads(2)
    {
        // Thread 0 :
        if (omp_get_thread_num() == 0)
            a = 2;
        // Thread 1 :
        if (omp_get_thread_num() == 1)
            b = 2;

        #pragma omp flush(a,b)
        printf("Thread %d: a=%d,b=%d\n", omp_get_thread_num(), a, b );
    }
    ...
}
```