

MPI


ECSE 420 - Tutorial 5
Dimitrios Stamoulis

TR 4110
October 22, 2014

Lab 3 – Introduction

Message Passing Architectures

- Complete computer as building block, including I/O
 - Communication via explicit I/O operations
- Programming model
 - direct access only to private address space (local memory),
 - communication via explicit messages (send/receive)
- High-level block diagram
 - Communication integration?
 - Mem, I/O, LAN, Cluster
 - Easier to build and scale than SAS
- Programming model more removed from basic hardware operations
 - Library or OS intervention



From ECSE 420
lecture slides!!

Software


Boundary

Hardware

During Lab 2...

We discussed about

- Shared memory architectures
- Shared address space programming model (OpenMP)

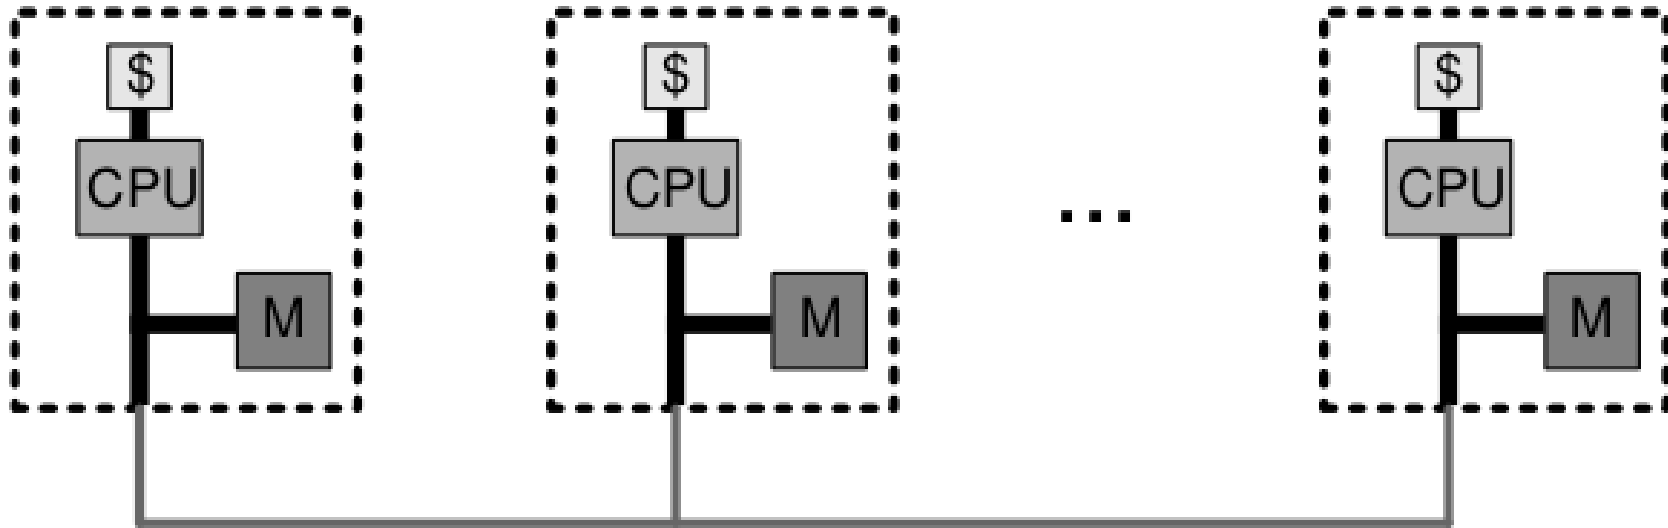


++ popular
architectural approach

++ simple
programming model

++ efficient parallel
programming

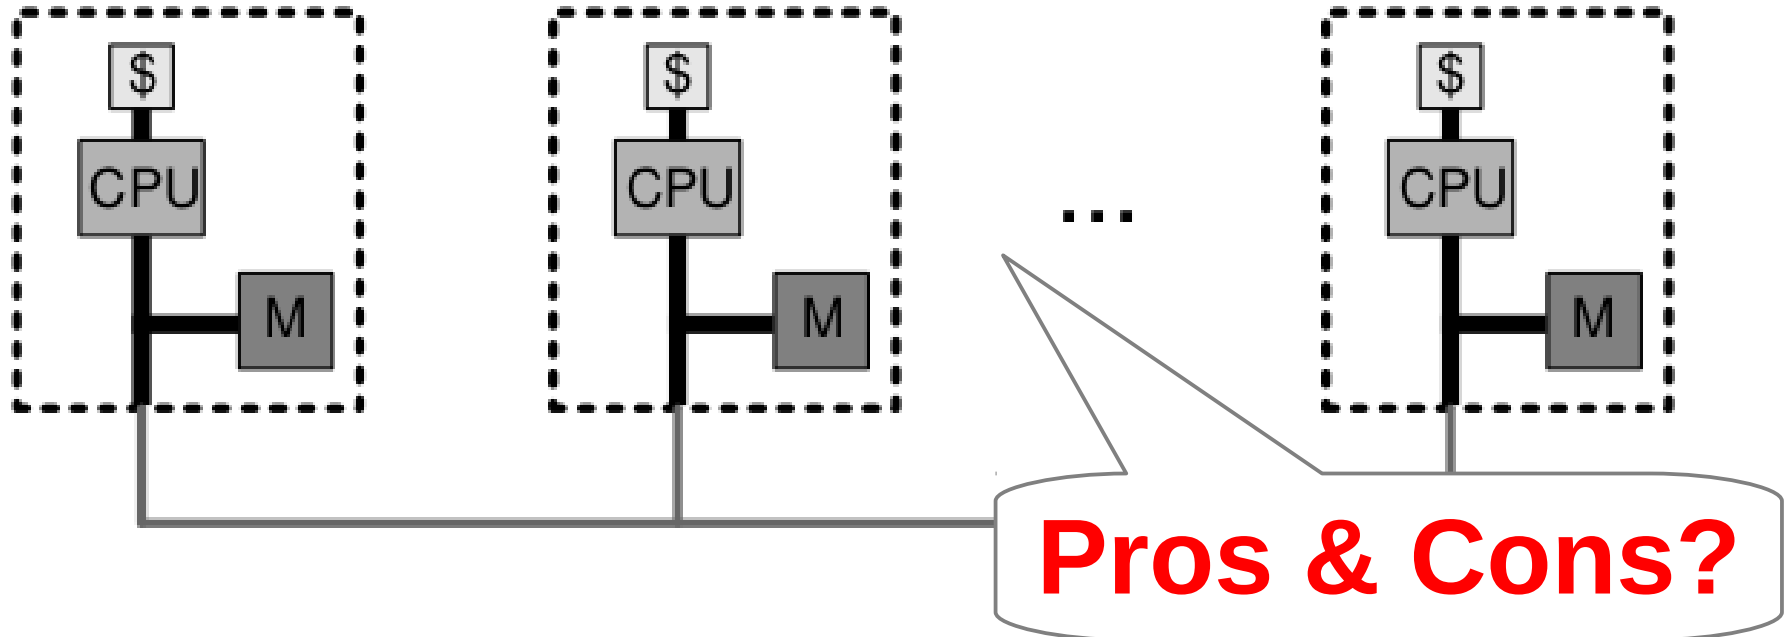
Message Passing Architectures



Programming model : we have to change it properly

→ More removed from basic hardware operations

Message Passing Architectures




Programming model : we have to change it properly

→ More removed from basic hardware operations

Send & receive operations (1/2)

We have:

- `send(void *sendbuf, int nelems, int dest)`
- `receive(void *recvbuf, int nelems, int source)`



Send & receive operations (2/2)

We have:

- `send(void *sendbuf, int nelems, int dest)`
- `receive(void *recvbuf, int nelems, int source)`

P0 sends data to P1

P0	P1
<code>a = 100;</code>	<code>receive(&a, 1, 0)</code>
<code>send(&a, 1, 1);</code>	<code>printf("%d\n", a);</code>
<code>a = b;</code>	

Good semantics → P1 receives 100

Bad semantics → P1 receives b

Send & receive operations (2/2)

We have:

- `send(void *sendbuf, int nelems, int dest)`
- `receive(void *recvbuf, int nelems, int source)`

P0 sends data to P1

P0

```
a = 100;  
send(&a, 1, 1);  
a = b;
```

P1

```
receive(&a, 1, 0)  
printf("%d\n", a);
```

Good semantics → P1 receives 100

Bad semantics → P1 receives b



HOW?

Programmer's challenges - Deadlock

We assume blocking, non-buffered send/receive:

P0

```
send(&b, 1, 1);  
receive(&a, 1, 1);  
0);
```

P1

```
send(&b, 1, 0);  
receive(&a, 1,
```

Both sends wait for both receives :

 **DEADLOCK**

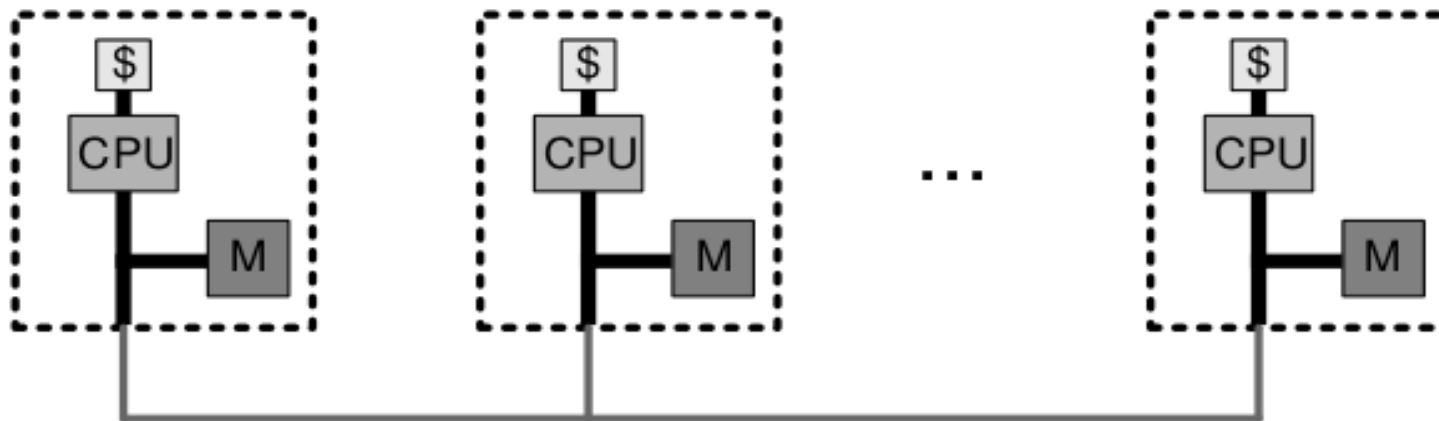
WHY?

Programmer's challenges

- Structure of message passing programs :
SPMD (single program, multiple data) model
 - All processes execute the same code
 - The process' execution flow differs based on its rank
- Programmer's challenges :
 - Maximum parallelization
 - Efficient usage of HW resources (e.g memory)
 - Minimum data to be transferred
 - Minimum number of messages
 - Minimum synchronization effort

Message Passing Programming Model

Programming model : we have to change it properly
→ More removed from basic hardware operations



- Message Passing Libraries
 - Vendors all had their own message passing libraries
- MPI :
 - Defines syntax, semantics of core set of library routines

Core set of routines for MPI

<code>MPI_Init</code>	Initializes MPI.
<code>MPI_Finalize</code>	Terminates MPI.
<code>MPI_Comm_size</code>	Determines the number of processes.
<code>MPI_Comm_rank</code>	Determines the label of calling process.
<code>MPI_Send</code>	Sends a message.
<code>MPI_Recv</code>	Receives a message.

Starting and Terminating MPI

```
int MPI_Init(int *argc, char ***argv)
```

- MPI_Init is called prior to other MPI routines-it initializes the MPI environment

```
int MPI_Finalize()
```

- MPI_Finalize is called at the end-it does clean-up

Return code for both is MPI_success

Communicators

Communicators :

- Are variables of type MPI_Comm.
- Define a set of processes that can communicate with each other

MPI_COMM_WORLD

- default communicator, all processes in program

```
int MPI_Comm_size(MPI_Comm comm, int *size)
```

- MPI_Comm_size - number of processes in communicator

```
int MPI_Comm_rank(MPI_Comm comm, int *rank)
```

Rank id's each process

Typical structure of MPI code

```
#include <mpi.h>

main(int argc, char *argv[])
{
    ...
    MPI_Init(&argc, &argv);
    MPI_Comm_size(MPI_COMM_WORLD, &size);
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);
    ...
    MPI_Finalize();
}
```

Our first example

```
#include <mpi.h>

main(int argc, char *argv[])
{
    int npes, myrank;
    MPI_Init(&argc, &argv);
    MPI_Comm_size(MPI_COMM_WORLD, &npes);
    MPI_Comm_rank(MPI_COMM_WORLD, &myrank);
    printf("I'm process %d out of %d\n", myrank, npes);
    MPI_Finalize();
}
```

- Compilation: \$ mpicc.mpich ex1.c -o ex1
- Execution: \$ mpiexec.mpich -np 2 ./ex1

MPI_Send()

```
int MPI_Send(void *buf, int count, MPI_Datatype  
             datatype, int dest, int tag, MPI_Comm comm)
```

- MPI_Send sends data in buf
- dest = rank of destination process
- Length of message = number of entries of type MPI_Datatype
- tag = type of message (e.g. a number)
- count & datatype specify length of buffer

Example :

```
int message[20], dest=1, tag=55;  
MPI_Send(message, 20, MPI_INT, dest, tag,  
          MPI_COMM_WORLD);
```

MPI_Recv()

```
int MPI_Recv(void *buf, int count, MPI_Datatype
             datatype, int source, int tag,
             MPI_Comm comm, MPI_Status *status)
```

- buf is where received message is stored
- Message to be received from source process
- status variable used to get info on MPI_Recv status

Example :

```
int message[50], source=0, tag=55;
MPI_Status status;
MPI_Recv(message, 50, MPI_INT, source, tag,
          MPI_COMM_WORLD, &status);
```

Datatypes

MPI Datatype

MPI_CHAR
MPI_SHORT
MPI_INT
MPI_LONG
MPI_UNSIGNED_CHAR
MPI_UNSIGNED_SHORT
MPI_UNSIGNED
MPI_UNSIGNED_LONG
MPI_FLOAT
MPI_DOUBLE
MPI_LONG_DOUBLE
MPI_BYTE

C Datatype

signed char
signed short int
signed int
signed long int
unsigned char
unsigned short int
unsigned int
unsigned long int
float
double
long double


Let's see a simple example

```
//Compute f(0) and f(1) in parallel and find their sum
#include <mpi.h>

int main(int argc, char** argv) {

    int v0, v1, sum, rank;
    MPI_Status status;
    MPI_Init(&argc, &argv);
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);
    if (rank == 1) {
        v1 = f(1);
        MPI_Send(&v1, 1, MPI_INT, 0, 50, MPI_COMM_WORLD);
    }
    else if (rank == 0) {
        v0 = f(0);
        MPI_Recv(&v1, 1, MPI_INT, 1, 50, MPI_COMM_WORLD, &status);
        sum = v0 + v1;
        printf("f(0)+f(1) = %d + %d = %d\n", v0, v1, sum);
    }
    MPI_Finalize();
}
```

Communication (1/2)



Communication (2/2)

We can have different types of communication :

→ Synchronous (MPI_Ssend) vs Buffered (MPI_Bsend)

What I consider to be MPI_Send 's successful completion?

→ Blocking vs Non-blocking (MPI_Isend)

Shall I wait for MPI_Recv to be completed?

→ Point-to-Point vs Collective

Can I send a message to everyone else?

Sending/Receiving Messages

```
int MPI_Recv(void *buf, int count, MPI_Datatype
             datatype, int source, int tag,
             MPI_Comm comm, MPI_Status *status)
```

→ A blocking operation : Returns only after message is in buffer..!!

```
int MPI_Send(void *buf, int count, MPI_Datatype
             datatype, int dest, int tag, MPI_Comm comm)
```

→ Synchronous : Returns after MPI_Recv issued and message is sent

→ Buffering : Returns after MPI_Send copied message into buffer
→ Does **NOT** wait for MPI_Recv to be issued

Deadlocks

```
int a[10], b[10], rank;
MPI_Status status;
...
MPI_Comm_rank(MPI_COMM_WORLD, &rank);
if (rank == 0) {
    MPI_Send(a, 10, MPI_INT, 1, 1, MPI_COMM_WORLD);
    MPI_Send(b, 10, MPI_INT, 1, 2, MPI_COMM_WORLD);
}
else if (rank == 1) {
    MPI_Recv(b, 10, MPI_INT, 0, 2, MPI_COMM_WORLD);
    MPI_Recv(a, 10, MPI_INT, 0, 1, MPI_COMM_WORLD);
}
```


Deadlocks

```
int a[10], b[10], rank;
MPI_Status status;
...
MPI_Comm_rank(MPI_COMM_WORLD, &rank);
if (rank == 0) {
    MPI_Send(a, 10, MPI_INT, 1, 1, MPI_COMM_WORLD);
    MPI_Send(b, 10, MPI_INT, 1, 2, MPI_COMM_WORLD);
}
else if (rank == 1) {
    MPI_Recv(a, 10, MPI_INT, 0, 2, MPI_COMM_WORLD);
    MPI_Recv(b, 10, MPI_INT, 0, 1, MPI_COMM_WORLD);
}
```


With a synchronous Send, we have a deadlock
e.g let's try MPI_Ssend()

MPI_Bcast - Collective Comm (1/2)

We could do something like :

```
...  
if (rank==0) {  
    for (dest=1; dest<size; dest++)  
        MPI_Send(msg, count, dest, tag, MPI_FLOAT, MPI_COMM_WORLD) ;  
}
```

...




Is that efficient? → We need **size-1** communication steps!!

MPI_Bcast - Collective Comm (2/2)

```
int MPI_Bcast(void* message, int count, MPI_Datatype  
              datatype, int root, MPI_Comm comm)
```

→ message is sent from process rank to all other processes of the communicator comm



What about now? Efficient? → We need **$\log_2(\text{size})$** communication steps!!

MPI_Barrier & MPI_Reduce

```
int MPI_Barrier(MPI_Comm comm)
```

→ Call returns after all processes have called the function

```
int MPI_Reduce(void *sendbuf, void *recvbuf, int count,  
               MPI_Datatype datatype, MPI_Op op, int  
               target, MPI_Comm comm)
```

→ Combines the `sendbuf` operand of each process using the operation specified by `MPI_Op op`.

→ Returns combined values in `recvbuf` of process with rank `target`

MPI_Barrier & MPI_Bcast - Example

```
int rank;
MPI_Init(&argc, &argv);
MPI_Status status;
int key = 0;

...
MPI_Comm_rank(MPI_COMM_WORLD, &rank);
if (rank == 0) {
    key = 1;
}
MPI_Bcast(&key, 1, MPI_INT, 0, MPI_COMM_WORLD);
MPI_Barrier(MPI_COMM_WORLD);

if (rank != 0) {
    printf("I am %d and I got the key=%d\n", rank, key);
}

MPI_Finalize();
```

Reduction Types

MPI_MAX	Maximum	C integers and floating point
MPI_MIN	Minimum	C integers and floating point
MPI_SUM	Sum	C integers and floating point
MPI_PROD	Product	C integers and floating point
MPI_LAND	Logical AND	C integers
MPI_BAND	Bit-wise AND	C integers and byte
MPI_LOR	Logical OR	C integers
MPI_BOR	Bit-wise OR	C integers and byte
MPI_LXOR	Logical XOR	C integers
MPI_BXOR	Bit-wise XOR	C integers and byte
...		

MPI_Reduce example

```
//Compute f(0) and f(1) in parallel and find their sum
#include <mpi.h>
```

```
int main(int argc, char** argv) {

    int mypart, sum, rank;
    MPI_Status status;
    MPI_Init(&argc, &argv);
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);
    mypart = f(rank);
    MPI_Reduce(&mypart, &sum, 1, MPI_INT, MPI_SUM, 0,
               MPI_COMM_WORLD);

    MPI_Finalize();
}
```